

```
import numpy as np
import pandas as pd

import matplotlib.pyplot as plt
import seaborn as sns

from datetime import datetime

from sklearn.preprocessing import StandardScaler
from sklearn.impute import SimpleImputer
imputer = SimpleImputer(missing_values=np.nan, strategy='mean')
from sklearn.model_selection import GridSearchCV, train_test_split, cross_val_score
from sklearn.metrics import classification_report, confusion_matrix
from sklearn.linear_model import LogisticRegression
from sklearn.naive_bayes import GaussianNB
from sklearn.neighbors import KNeighborsClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.ensemble import GradientBoostingClassifier
from sklearn.svm import SVC
from sklearn.metrics import roc_curve, auc
import os
import warnings
warnings.filterwarnings('ignore')

data=pd.read_csv('/content/heart.csv')

#Now, I will check null on all data and If data has null, I will sum of null data's. In this w
print('Data Sum of Null Values\n')
```

Saved successfully!

```
Data Sum of Null Values

age          0
sex          0
cp           0
trestbps     0
chol         0
fbs          0
restecg      0
thalach      0
exang        0
oldpeak      0
slope        0
ca           0
thal         0
target       0
dtype: int64
```


```
#all rows control for null values
print(data.isnull().values.any())
```

```
False
```

```
# TODO:- Age Analysis
print(data.age.value_counts()[:10])
#data age show value counts for age least 10
```

```
58    19
57    17
54    16
59    14
52    13
51    12
62    11
44    11
60    11
56    11
Name: age, dtype: int64
```

```
sns.barplot(x=data.age.value_counts()[:10].index,y=data.age.value_counts()[:10].values)
plt.xlabel('Age')
plt.ylabel('Age Counter')
plt.title('Different age group')
plt.show()
```



```
#firstly find min and max ages
minage=min(data.age)
maxage=max(data.age)
meanage=data.age.mean()
print('Min Age :',minage)
print('Max Age :',maxage)
print('Mean Age :',meanage)
```

```

Min Age : 29
Max Age : 77
Mean Age : 54.366336633663366

```

```

young_ages=data[(data.age>=29)&(data.age<40)]
middle_ages=data[(data.age>=40)&(data.age<55)]
elderly_ages=data[(data.age>55)]
print('Young Ages :',len(young_ages))
print('Middle Ages :',len(middle_ages))
print('Elderly Ages :',len(elderly_ages))

```

```

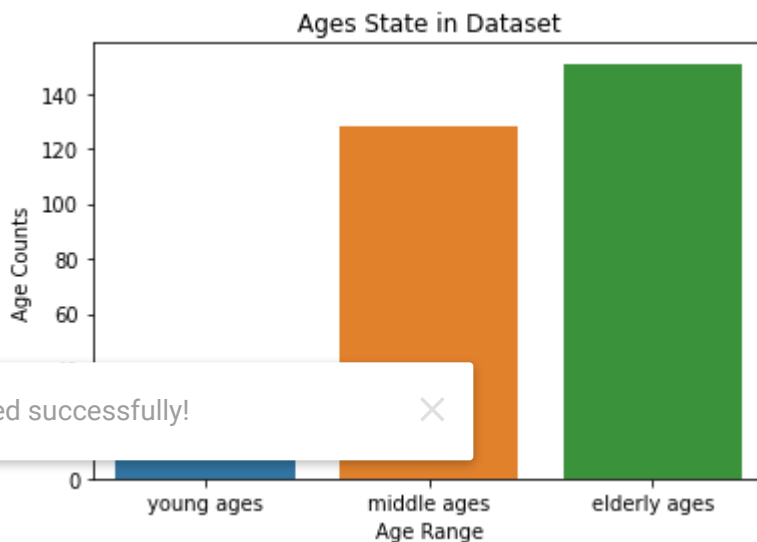
Young Ages : 16
Middle Ages : 128
Elderly Ages : 151

```

```

sns.barplot(x=['young ages','middle ages','elderly ages'],y=[len(young_ages),len(middle_ages),len(elderly_ages)])
plt.xlabel('Age Range')
plt.ylabel('Age Counts')
plt.title('Ages State in Dataset')
plt.show()

```



```

data['ageRange']=0
youngage_index=data[(data.age>=29)&(data.age<40)].index
middleage_index=data[(data.age>=40)&(data.age<55)].index
elderlyage_index=data[(data.age>55)].index

```

```

for index in elderlyage_index:
    data.loc[index,'ageRange']=2

```

```

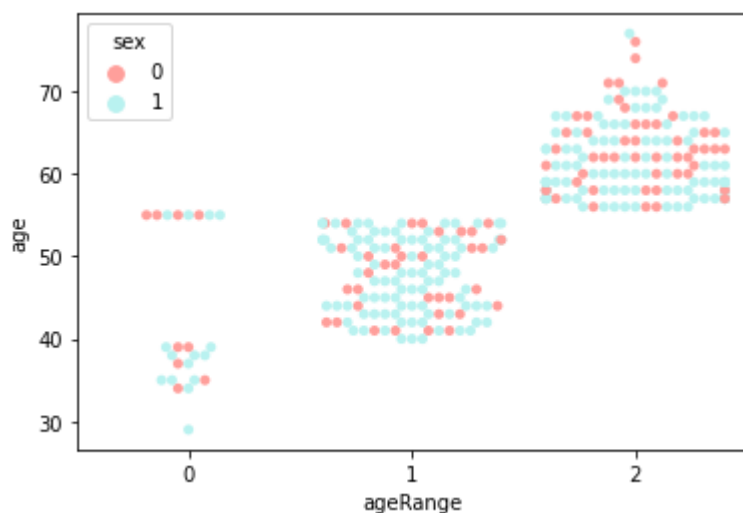
for index in middleage_index:
    data.loc[index,'ageRange']=1

```

```

for index in youngage_index:
    # Draw a categorical scatterplot to show each observation
    sns.swarmplot(x="ageRange", y="age", hue='sex',
                  palette=["r", "c", "y"], data=data)
plt.show()


```



```

# Plot the total crashes
sns.set_color_codes("pastel")
sns.barplot(y="ageRange", x="sex", data=data,
            label="Total", color="b")
plt.show()

```




```
print(elderly_ages.groupby(elderly_ages['sex'])['thalach'].agg('sum'))
```

```


sex
0    7578
1   13948
Name: thalach, dtype: int64

```

```
sns.barplot(x=elderly_ages.groupby(elderly_ages['sex'])['thalach'].agg('sum').index,y=elderly_
plt.title("Gender Group Thalach Show Sum Time")
plt.show()
```



```
colors = ['blue','green','yellow']
explode = [0,0,0.1]
plt.figure(figsize = (5,5))
#plt.pie([target_0_agerang_0,target_1_agerang_0], explode=explode, labels=['Target 0 Age Rang
plt.pie([len(young_ages),len(middle_ages),len(elderly_ages)],labels=['young ages','middle age
plt.title('Age States',color = 'blue',fontsize = 15)
plt.show()
```




```
# TODO :- Sex (Gender) Analysis
```

```
print(data.sex.value_counts())
```

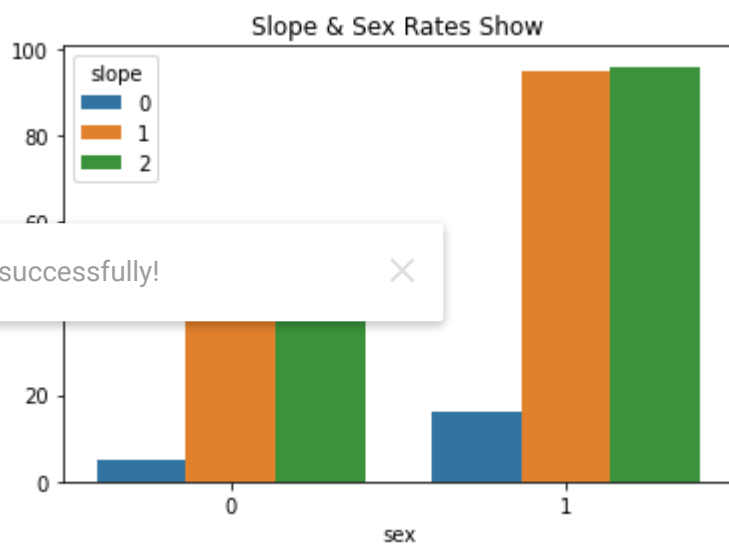
```
1    207
0     96
```

Name: sex, dtype: int64

```
#Sex (1 = male; 0 = female)
sns.countplot(data.sex)
plt.show()
```



```
sns.countplot(data.sex, hue=data.slope)
plt.title('Slope & Sex Rates Show')
plt.show()
```



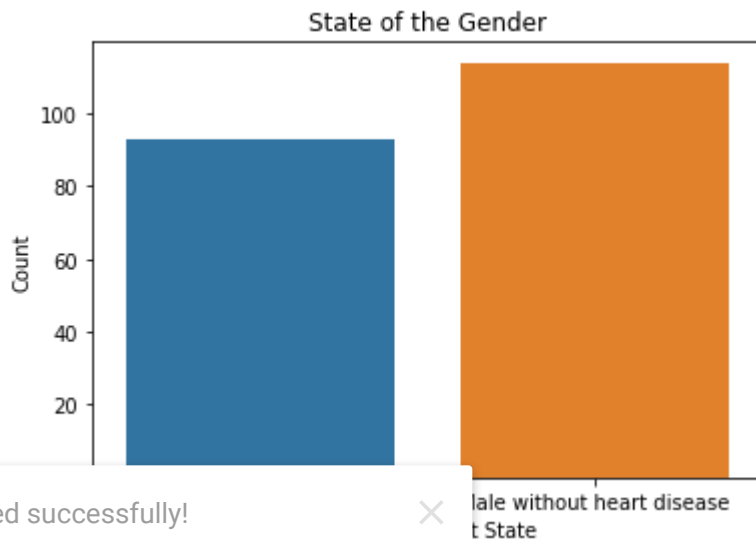
```
total_genders_count=len(data.sex)
male_count=len(data[data['sex']==1])
female_count=len(data[data['sex']==0])
print('Total Genders :',total_genders_count)
print('Male Count    :',male_count)
print('Female Count  :',female_count)
```

```
Total Genders : 303
Male Count    : 207
Female Count  : 96
```

```
#Percentage ratios
print("Male State: {:.2f}%".format((male_count / (total_genders_count)*100)))
print("Female State: {:.2f}%".format((female_count / (total_genders_count)*100)))
```


```
Male State: 68.32%
Female State: 31.68%
```

```
# Male State & target 1 & 0
male_andtarget_on=len(data[(data.sex==1)&(data['target']==1)])
male_andtarget_off=len(data[(data.sex==1)&(data['target']==0)])
####
sns.barplot(x=['Male with heart disease','Male without heart disease'],y=[male_andtarget_on,male_andtarget_off])
plt.xlabel('Male and Target State')
plt.ylabel('Count')
plt.title('State of the Gender')
plt.show()
```



Saved successfully!

```
#Female State & target 1 & 0
female_andtarget_on=len(data[(data.sex==0)&(data['target']==1)])
female_andtarget_off=len(data[(data.sex==0)&(data['target']==0)])
####
sns.barplot(x=['Female with heart disease','Female without heart disease'],y=[female_andtarget_on,female_andtarget_off])
plt.xlabel('Female and Target State')
plt.ylabel('Count')
plt.title('State of the Gender')
plt.show()
```

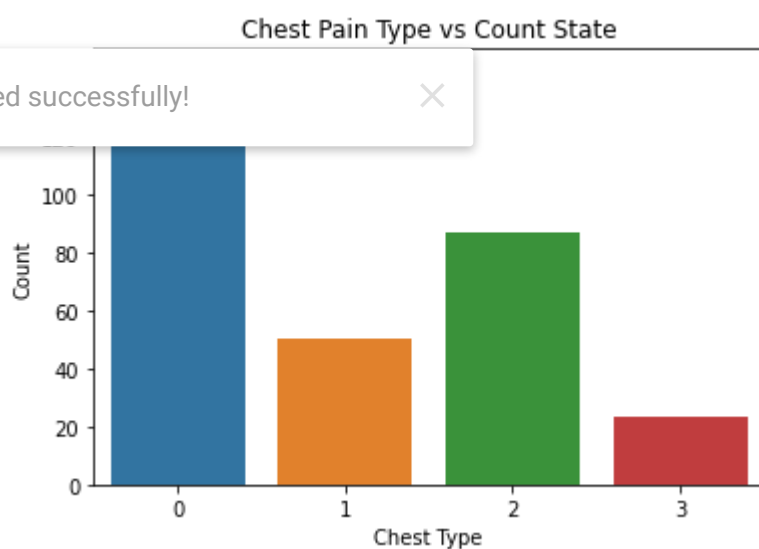


# TODO:- Chest Pain Type Analysis

```
print(data.cp.value_counts())
```


```
0    143
2     87
1     50
3     23
Name: cp, dtype: int64
```

```
sns.countplot(data.cp)
plt.xlabel('Chest Type')
plt.ylabel('Count')
plt.title('Chest Pain Type vs Count State')
plt.show()
#0 status at least
#1 condition slightly distressed
#2 condition medium problem
#3 condition too bad
```



```
# Show the results of a linear regression within each dataset
sns.lmplot(x="trestbps", y="chol", data=data, hue="cp")
plt.show()
```





```
data.describe(include=[np.number])
```

	age	sex	cp	trestbps	chol	fbs	restecg
<b>count</b>	303.000000	303.000000	303.000000	303.000000	303.000000	303.000000	303.000000
<b>mean</b>	54.366337	0.683168	0.966997	131.623762	246.264026	0.148515	0.528053
<b>std</b>	9.082101	0.466011	1.032052	17.538143	51.830751	0.356198	0.525860
<b>min</b>	29.000000	0.000000	0.000000	94.000000	126.000000	0.000000	0.000000
<b>25%</b>	47.500000	0.000000	0.000000	120.000000	211.000000	0.000000	0.000000
			1.000000	130.000000	240.000000	0.000000	1.000000
			2.000000	140.000000	274.500000	0.000000	1.000000
<b>max</b>	77.000000	1.000000	3.000000	200.000000	564.000000	1.000000	2.000000

Saved successfully!



```
dt= data.copy()
dt.head()
```



	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	ca	thal
<b>0</b>	63	1	3	145	233	1	0	150	0	2.3	0	0	1
<b>1</b>	37	1	2	130	250	0	1	187	0	3.5	0	0	2
<b>2</b>	41	0	1	130	204	0	0	172	0	1.4	2	0	2
<b>3</b>	56	1	1	120	236	0	1	178	0	0.8	2	0	2
<b>4</b>	57	0	0	120	354	0	1	163	1	0.6	2	0	2

```
dt=dt.rename(columns={'cp':'chest_pain_type','restecg':'rest_ecg','slope':'st_slope'})

# Plotting attrition of employees
fig, (ax1, ax2) = plt.subplots(nrows=1, ncols=2, sharey=False, figsize=(14,6))

ax1 = dt['target'].value_counts().plot.pie( x="Heart disease" ,y = 'no.of patients',
      autopct = "%1.0f%%",labels=["Heart Disease","Normal"], startangle = 60,ax=
ax1.set(title = 'Percentage of Heart disease patients in Dataset')

ax2 = dt["target"].value_counts().plot(kind="barh" ,ax =ax2)
for i,j in enumerate(dt["target"].value_counts().values):
    ax2.text(.5,i,j,fontsize=12)
ax2.set(title = 'No. of Heart disease patients in Dataset')
plt.show()
```





Saved successfully!



```
plt.figure(figsize=(18,12))
plt.subplot(221)
dt["sex"].value_counts().plot.pie(autopct = "%1.0f%",colors = sns.color_palette("prism",5),s
wedgeprops={"linewidth":2,"edgecolor":"k"},explode=[.1,.1],shadow =True)
plt.title("Distribution of Gender")
plt.subplot(222)
ax= sns.distplot(dt['age'], rug=True)
plt.title("Age wise distribution")
plt.show()
```





```
# creating separate df for normal and heart patients
```

```
attr_1=dt[dt['target']==1]
```

```
attr_0=dt[dt['target']==0]
```

```
# plotting normal patients
```

```
fig = plt.figure(figsize=(15,5))
```

```
ax1 = plt.subplot2grid((1,2),(0,0))
```

```
sns.distplot(attr_0['age'])
```

```
plt.title('AGE DISTRIBUTION OF NORMAL PATIENTS', fontsize=15, weight='bold')
```

```
ax1 = plt.subplot2grid((1,2),(0,1))
```

```
palette='viridis')
```

```
plt.title('GENDER DISTRIBUTION OF NORMAL PATIENTS', fontsize=15, weight='bold' )
```

Saved successfully!

```
#plotting heart patients
```

```
fig = plt.figure(figsize=(15,5))
```

```
ax1 = plt.subplot2grid((1,2),(0,0))
```

```
sns.distplot(attr_1['age'])
```





```
plt.title('AGE DISTRIBUTION OF HEART DISEASE PATIENTS', fontsize=15, weight='bold')
```

```
ax1 = plt.subplot2grid((1,2),(0,1))
```

```
sns.countplot(attr_1['sex'], palette='viridis')
```

```
plt.title('GENDER DISTRIBUTION OF HEART DISEASE PATIENTS', fontsize=15, weight='bold' )
```

```
plt.show()
```



Saved successfully!



```
dt.chest_pain_type.unique()
```

```
array([3, 2, 1, 0])
```

```
print(dt.chest_pain_type.unique())
```

```
print(dt.rest_ecg.unique())
```

```
print(dt.st_slope.unique())
```

```
[3 2 1 0]
```

```
[0 1 2]
```

```
[0 2 1]
```

```
# converting features to categorical features
```

```
dt['chest_pain_type'][dt['chest_pain_type'] == 0] = 'typical angina'
```

```
dt['chest_pain_type'][dt['chest_pain_type'] == 1] = 'atypical angina'
```

```
dt['chest_pain_type'][dt['chest_pain_type'] == 2] = 'non-anginal pain'
```

```
dt['chest_pain_type'][dt['chest_pain_type'] == 3] = 'asymptomatic'
```

```
dt['rest_ecg'][dt['rest_ecg'] == 0] = 'normal'
```

```
dt['rest_ecg'][dt['rest_ecg'] == 1] = 'ST-T wave abnormality'
```

```
dt['rest_ecg'][dt['rest_ecg'] == 2] = 'left ventricular hypertrophy'
```

```
dt['st_slope'][dt['st_slope'] == 0] = 'upsloping'
```

```
dt['st_slope'][dt['st_slope'] == 1] = 'flat'
```

```
dt['st_slope'][dt['st_slope'] == 2] = 'downsloping'
```

```
dt["sex"] = dt.sex.apply(lambda x:'male' if x==1 else 'female')
```

```
dt.chest_pain_type.unique()
```

```
array(['asymptomatic', 'non-anginal pain', 'atypical angina',  
      'typical angina'], dtype=object)
```

```
dt.head()
```

	age	sex	chest_pain_type	trestbps	chol	fbs	rest_ecg	thalach	exang	oldpeak
0	63	male	asymptomatic	145	233	1	normal	150	0	2.3
1	37	male	non-anginal pain	130	250	0	ST-T wave abnormality	187	0	3.5
2	41	female	atypical angina	130	204	0	normal	172	0	1.4
3	56	male	typical angina	120	236	0	ST-T wave abnormality	178	0	0.8

Saved successfully!

```
dt['rest_ecg'].value_counts()
```

```
ST-T wave abnormality    152
normal                   147
left ventricular hypertrophy    4
Name: rest_ecg, dtype: int64
```

```
dt['chest_pain_type'].value_counts()
```

```

typical angina      143
non-anginal pain    87
atypical angina     50
asymptomatic        23
Name: chest_pain_type, dtype: int64

```

```
dt['st_slope'].value_counts()
```

```

downsloping    142
flat           140
upsloping       21
Name: st_slope, dtype: int64

```

```
dt.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 303 entries, 0 to 302
Data columns (total 15 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   age                   303 non-null   int64
 1   sex                   303 non-null   object
 2   chest_pain_type       303 non-null   object
 3   trestbps              303 non-null   int64
 4   chol                  303 non-null   int64
 5   fbs                   303 non-null   int64
 6   rest_ecg              303 non-null   object
 7   thalach               303 non-null   int64
 8   exang                 303 non-null   int64
 9   oldpeak               303 non-null   float64
10   st_slope              303 non-null   object
11   ageRange              303 non-null   int64
12   max_heart_rate        303 non-null   int64
13   max_heart_rate_time    303 non-null   int64
14   ageRange              303 non-null   int64
dtypes: float64(1), int64(10), object(4)
memory usage: 35.6+ KB

```



Saved successfully!

```

# plotting normal patients
fig = plt.figure(figsize=(15,5))
ax1 = plt.subplot2grid((1,2),(0,0))
sns.countplot(attr_0['chest_pain_type'])
plt.title('CHEST PAIN OF NORMAL PATIENTS', fontsize=15, weight='bold')

#plotting heart patients
ax1 = plt.subplot2grid((1,2),(0,1))
sns.countplot(attr_1['chest_pain_type'], palette='viridis')
plt.title('CHEST PAIN OF HEART PATIENTS', fontsize=15, weight='bold' )
plt.show()

```



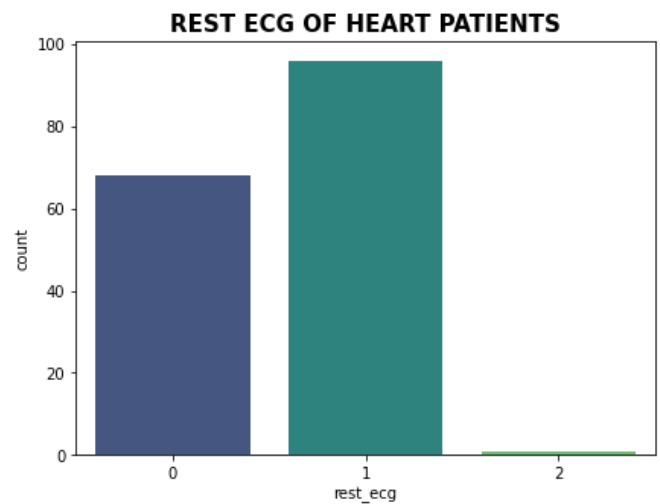

```
#Exploring the Heart Disease patients based on Chest Pain Type
plot_criteria= ['chest_pain_type', 'target']
cm = sns.light_palette("red", as_cmap=True)
(round(pd.crosstab(dt[plot_criteria[0]], dt[plot_criteria[1]], normalize='columns') * 100,2))
```

	target	0	1
chest_pain_type			
asymptomatic		5.070000	9.700000
atypical angina		6.520000	24.850000
typical angina		7.000000	20.640000

Saved successfully!

```
# plotting normal patients
fig = plt.figure(figsize=(15,5))
ax1 = plt.subplot2grid((1,2),(0,0))
sns.countplot(attr_0['rest_ecg'])
plt.title('REST ECG OF NORMAL PATIENTS', fontsize=15, weight='bold')

#plotting heart patients
ax1 = plt.subplot2grid((1,2),(0,1))
sns.countplot(attr_1['rest_ecg'], palette='viridis')
plt.title('REST ECG OF HEART PATIENTS', fontsize=15, weight='bold' )
plt.show()
```



```
#Exploring the Heart Disease patients based on REST ECG
plot_criteria= ['rest_ecg', 'target']
cm = sns.light_palette("red", as_cmap=True)
(round(pd.crosstab(dt[plot_criteria[0]], dt[plot_criteria[1]], normalize='columns') * 100,2))
```


	target	
	0	1
rest_ecg		
<b>ST-T wave abnormality</b>	40.580000	58.180000
<b>left ventricular hypertrophy</b>	2.170000	0.610000
<b>normal</b>	57.250000	41.210000

Saved successfully!

```
ax1 = plt.subplot2grid((1,2),(0,0))
sns.countplot(attr_0['st_slope'])
plt.title('ST SLOPE OF NORMAL PATIENTS', fontsize=15, weight='bold')
```

```
#plotting heart patients
ax1 = plt.subplot2grid((1,2),(0,1))
sns.countplot(attr_1['st_slope'], palette='viridis')
plt.title('ST SLOPE OF HEART PATIENTS', fontsize=15, weight='bold' )
plt.show()
```





```
# summary statistics of categorical columns
dt.describe(include =[np.object])
```

	sex	chest_pain_type	rest_ecg	st_slope
count	303	303	303	303
unique	2	4	3	3
top	male	typical angina	ST-T wave abnormality	downsloping
freq	207	143	152	142

```
# TODO:- Age Range Analysis
```

```
target_0_agerang_0=len(data[(data.target==0)&(data.ageRange==0)])
target_1_agerang_0=len(data[(data.target==1)&(data.ageRange==0)])
```

```
colors = ['blue','green']
```

Saved successfully!

```
plt.pie([target_0_agerang_0,target_1_agerang_0], explode=explode, labels=['Target 0 (without
plt.title('Target(patient with and without heart disease) vs Age Range (Young Age) ',color =
plt.show()
```


## Target(patient with and without heart disease) vs Age Range (Young Age)

```

target_0_agerang_1=len(data[(data.target==0)&(data.ageRange==1)])
target_1_agerang_1=len(data[(data.target==1)&(data.ageRange==1)])
colors = ['blue','green']
explode = [0.1,0]
plt.figure(figsize = (5,5))
plt.pie([target_0_agerang_1,target_1_agerang_1], explode=explode, labels=['Target 0 (without heart disease) Age Range 0 (young age)', 'Target 1 (with heart disease) Age Range 0 (young age)'], color = colors, autopct='%1.1f%%')
plt.title('Target (patient with and without heart disease) vs Age Range (Middle Age)',color = 'red')
plt.show()

```

## Target (patient with and without heart disease) vs Age Range (Middle Age)



```

target_0_agerang_2=len(data[(data.target==0)&(data.ageRange==2)])
target_1_agerang_2=len(data[(data.target==1)&(data.ageRange==2)])

```

Saved successfully!

```

explode = [0,0.1]
plt.figure(figsize = (5,5))
plt.pie([target_0_agerang_2,target_1_agerang_2], explode=explode, labels=['Target 0 (without heart disease) Age Range 2 (elderly age)', 'Target 1 (with heart disease) Age Range 2 (elderly age)'], color = colors, autopct='%1.1f%%')
plt.title('Target (patient with and without heart disease) vs Age Range (Elderly Age) ',color = 'red')
plt.show()

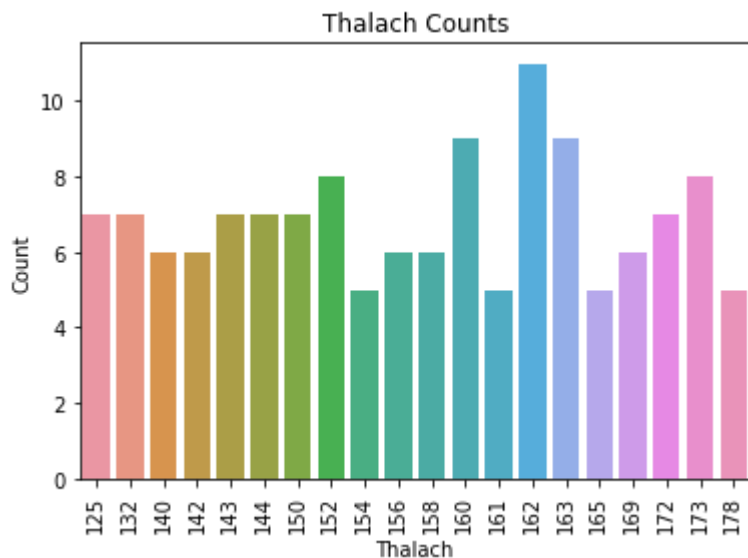
```

## Target (patient with and without heart disease) vs Age Range (Elderly Age)

Target 0 (without heart disease) Age Range 2 (elderly age)



```
sns.barplot(x=data.thalach.value_counts()[ :20].index,y=data.thalach.value_counts()[ :20].value)
plt.xlabel('Thalach')
plt.ylabel('Count')
plt.title('Thalach Counts')
plt.xticks(rotation=90)
plt.show()
```



```
age_range_thalach=data.groupby('ageRange')['thalach'].mean()
sns.barplot(x=age_range_thalach.index,y=age_range_thalach.values)
```

Saved successfully!



age')

ch to the age range')

```
plt.show()
```


```
#As shown in this graph, this rate decreases as the heart rate
#is faster and in old age areas.
```

illustration of the thalach to the age range



```
cp_thalach=data.groupby('cp')['thalach'].mean()
sns.barplot(x=cp_thalach.index,y=cp_thalach.values)
plt.xlabel('Degree of Chest Pain (Cp)')
plt.ylabel('Maximum Thalach By Cp Values')
plt.title('Illustration of thalach to degree of chest pain')
plt.show()
```

#As seen in this graph, it is seen that the heart rate is less  
 #when the chest pain is low. But in cases where chest pain is  
 #1, it is observed that the area is more. 2 and 3 were found to  
 #be of the same degree.



# TODO :- Thal Analysis

Saved successfully!

```
2    166
3    117
1     18
0      2
Name: thal, dtype: int64
```

```
#Target 1
a=len(data[(data['target']==1)&(data['thal']==0)])
b=len(data[(data['target']==1)&(data['thal']==1)])
c=len(data[(data['target']==1)&(data['thal']==2)])
d=len(data[(data['target']==1)&(data['thal']==3)])
print('Target 1 Thal 0: ',a)
print('Target 1 Thal 1: ',b)
print('Target 1 Thal 2: ',c)
print('Target 1 Thal 3: ',d)
```

#so,Apparently, there is a rate at Thal 2.Now, draw graph

```

print('*'*50)
#Target 0
e=len(data[(data['target']==0)&(data['thal']==0)])
f=len(data[(data['target']==0)&(data['thal']==1)])
g=len(data[(data['target']==0)&(data['thal']==2)])
h=len(data[(data['target']==0)&(data['thal']==3)])
print('Target 0 Thal 0: ',e)
print('Target 0 Thal 1: ',f)
print('Target 0 Thal 2: ',g)
print('Target 0 Thal 3: ',h)

Target 1 Thal 0: 1
Target 1 Thal 1: 6
Target 1 Thal 2: 130
Target 1 Thal 3: 28
*****
Target 0 Thal 0: 1
Target 0 Thal 1: 12
Target 0 Thal 2: 36
Target 0 Thal 3: 89

```

# TODO :- Target Analysis

```

print(data.target.unique())
#only two values are shown.
#A value of 1 is the value of patient 0

```

```
[1 0]
```

```

sns.countplot(data.target)
plt.xlabel('Target')

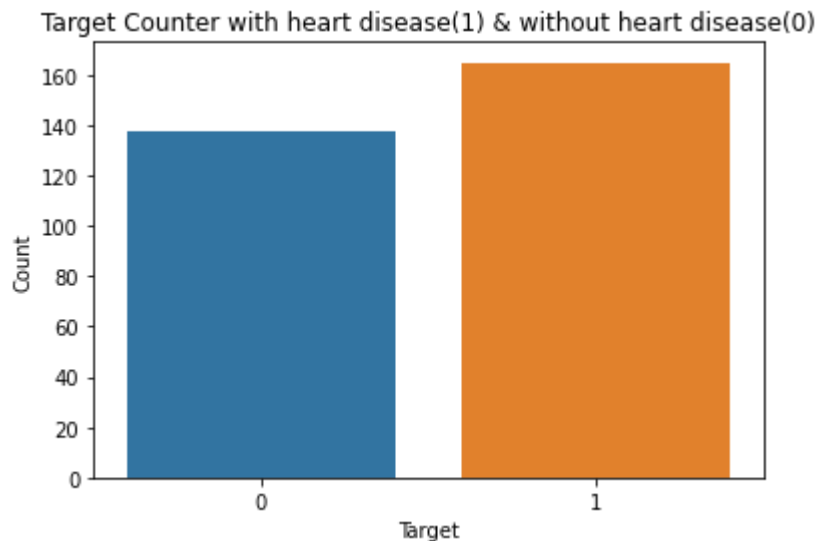
```

Saved successfully!

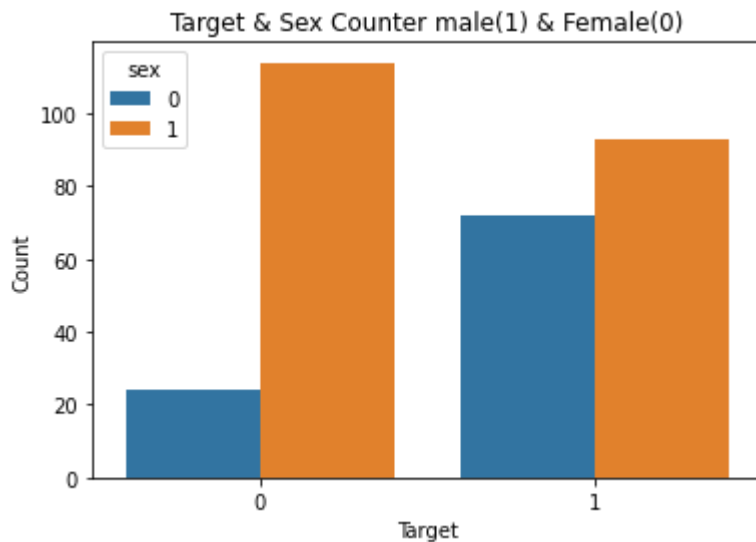


disease(1) & without heart disease(0)')

```
plt.show()
```



```
sns.countplot(data.target,hue=data.sex)
plt.xlabel('Target')
plt.ylabel('Count')
plt.title('Target & Sex Counter male(1) & Female(0)')
plt.show()
```




```
#determine the age ranges of patients with and without sickness and make analyzes about them
age_counter_target_1=[]
age_counter_target_0=[]
for age in data.age.unique():
    age_counter_target_1.append(len(data[(data['age']==age)&(data.target==1)]))
    age_counter_target_0.append(len(data[(data['age']==age)&(data.target==0)]))
```

```
#now draw show on graph
```

Saved successfully!

```
age_counter_target_1,color='blue',label='Patient with heart
plt.scatter(x=data.age.unique(),y=age_counter_target_0,color='red',label='Patient without heart
plt.legend(loc='upper right',frameon=True)
plt.xlabel('Age')
plt.ylabel('Count')
plt.title('Target 0 (Patient without heart disease) & Target 1 (Patient with heart disease) S
plt.show()
```

Target 0 (Patient without heart disease) &amp; Target 1 (Patient with heart disease) State



```
df=data.copy()
df.head()
```

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	ca	thal
0	63	1	3	145	233	1	0	150	0	2.3	0	0	1
1	37	1	2	130	250	0	1	187	0	3.5	0	0	2
2	41	0	1	130	204	0	0	172	0	1.4	2	0	2
3	56	1	1	120	236	0	1	178	0	0.8	2	0	2
4	57	0	0	120	354	0	1	163	1	0.6	2	0	2

```
from tensorflow.keras.utils import to_categorical
# one hot encode
encoded1 = pd.DataFrame(to_categorical(df.restecg),columns=['restecg0','restecg1','restecg2'])
encoded2 = pd.DataFrame(to_categorical(df.cp),columns=['cp0','cp1','cp2','cp3'])
encoded3 = pd.DataFrame(to_categorical(df.thal),columns=['thal0','thal1','thal2','thal3'])
encoded4 = pd.DataFrame(to_categorical(df.ca),columns=['ca0','ca1','ca2','ca3','ca4'])
encoded5 = pd.DataFrame(to_categorical(df.slope),columns=['slope0','slope1','slope2'])

df = pd.concat([df,encoded1,encoded2,encoded3,encoded4,encoded5],axis=1).drop(['restecg','cp']
df.head()
```

Saved successfully!



	age	sex	cp	trestbps	chol	fbs	thalach	exang	oldpeak	target	ageRange	restecg0	r
0	63	1	3	145	233	1	150	0	2.3	1	2	1.0	
1	37	1	2	130	250	0	187	0	3.5	1	0	0.0	
2	41	0	1	130	204	0	172	0	1.4	1	1	1.0	
3	56	1	1	120	236	0	178	0	0.8	1	2	0.0	
4	57	0	0	120	354	0	163	1	0.6	1	2	0.0	

```
X_train, X_test, y_train, y_test=train_test_split(
    df.drop(['target'], axis=1),
    df[['target']],
    test_size=0.3,
    random_state=41)
```

```

print(X_train.shape)
print(X_test.shape)

(212, 28)
(91, 28)

for column in X_train.columns:

    df_train1 = X_train[(y_train.target==0) & (X_train[column]<np.mean(X_train.loc[y_train.ta
    df_test1 = X_test[(y_test.target==0) & (X_test[column]<np.mean(X_train.loc[y_train.target

    label_train1 = y_train[(y_train.target==0) & (X_train[column]<np.mean(X_train.loc[y_train
    label_test1 = y_test[(y_test.target==0) & (X_test[column]<np.mean(X_train.loc[y_train.tar

    df_train2 = X_train[(y_train.target==1) & (X_train[column]<np.mean(X_train.loc[y_train.ta
    df_test2 = X_test[(y_test.target==1) & (X_test[column]<np.mean(X_train.loc[y_train.target

    label_train2 = y_train[(y_train.target==1) & (X_train[column]<np.mean(X_train.loc[y_train
    label_test2 = y_test[(y_test.target==1) & (X_test[column]<np.mean(X_train.loc[y_train.tar

X_train=pd.concat([df_train1,df_train2])
y_train=pd.concat([label_train1,label_train2])

X_test=pd.concat([df_test1,df_test2])
y_test=pd.concat([label_test1,label_test2])

from sklearn.feature_selection import VarianceThreshold
from sklearn.metrics import classification_report

Saved successfully! XGBClassifier

from sklearn import tree
from sklearn import metrics
constant_filter = VarianceThreshold(threshold=0.0)
constant_filter.fit(X_train)
X_train = constant_filter.transform(X_train)
X_test = constant_filter.transform(X_test)

X_train.shape, X_test.shape

((212, 28), (91, 28))

columns=df.columns
columns_new=[]
for i in columns:
    columns_new.append(any(df[i].isnull()|df[i].isnull()))
df=df.drop(columns[columns_new],axis=1)

```



```

mm_scaler = preprocessing.StandardScaler()
X_train = pd.DataFrame(mm_scaler.fit_transform(X_train))
X_test=pd.DataFrame(mm_scaler.transform(X_test))

from sklearn.linear_model import LogisticRegression

logreg = LogisticRegression(max_iter=50)
logreg.fit(X_train, y_train.values.ravel())
y_pred_logreg=logreg.predict(X_test)
acc = metrics.accuracy_score(y_pred_logreg,y_test.values.ravel())*100

print("Test Accuracy of Logistic Regression Algorithm: {:.2f}%".format(acc))

    Test Accuracy of Logistic Regression Algorithm: 84.62%

```

```

from sklearn.naive_bayes import GaussianNB
nb = GaussianNB()
nb.fit(X_train, y_train.values.ravel())

y_pred_nb=nb.predict(X_test)
acc = metrics.accuracy_score(y_pred_nb,y_test.values.ravel())*100

print("Accuracy of Naive Bayes: {:.2f}%".format(acc))

    Accuracy of Naive Bayes: 80.22%

```

```

from sklearn.svm import SVC
from sklearn import svm

```

Saved successfully!



```

probability=True)
l())

```

```

y_pred_svm=svm.predict(X_test)
acc = metrics.accuracy_score(y_pred_svm,y_test.values.ravel())*100

print("Test Accuracy of SVM Algorithm: {:.2f}%".format(acc))

    Test Accuracy of SVM Algorithm: 76.92%

```

```

# KNN Model
from sklearn.neighbors import KNeighborsClassifier

# try to find best k value
score = []


for i in range(1,20):
    knn = KNeighborsClassifier(n_neighbors = i) # n_neighbors means k
    knn.fit(X_train, y_train.values.ravel())

```

```
score.append(knn.score(X_test, y_test.values.ravel()))
```

```
plt.plot(range(1,20), score)
plt.xticks(np.arange(1,20,1))
plt.xlabel("K neighbors")
plt.ylabel("Score")
plt.show()
```

```
acc = max(score)*100
print("Maximum KNN Score is {:.2f}%".format(acc))
```



Maximum KNN Score is 89.01%

```
knn = KNeighborsClassifier(n_neighbors = 7) # n_neighbors means k
knn.fit(X_train, y_train.values.ravel())
```

```
KNeighborsClassifier(n_neighbors=7)
```

Saved successfully!

```
acc = metrics.accuracy_score(y_pred_knn,y_test.values.ravel())*100
print("KNN Accuracy Score : {:.2f}%".format(acc))
```

KNN Accuracy Score : 83.52%

```
# Random Forest Classification
from sklearn.ensemble import RandomForestClassifier
rf = RandomForestClassifier(class_weight="balanced",n_estimators=200,random_state = 1)
rf.fit(X_train, y_train.values.ravel())
y_pred_rf=rf.predict(X_test)
acc = metrics.accuracy_score(y_pred_rf,y_test.values.ravel())*100
print("Random Forest Algorithm Accuracy Score : {:.2f}%".format(acc))
```

Random Forest Algorithm Accuracy Score : 85.71%

```
from keras.models import Sequential
from keras.layers import Dense
```

```
from keras.layers import Dropout
```

```
# define the keras model
model = Sequential()
model.add(Dense(12, input_dim=X_train.shape[1], activation='relu'))
model.add(Dropout(0.2))
model.add(Dense(8, activation='relu'))
model.add(Dropout(0.2))
model.add(Dense(1, activation='sigmoid'))
# compile the keras model
model.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy'])
# fit the keras model on the dataset
model.fit(X_train, y_train, epochs=40, batch_size=32)
# evaluate the keras model
_, accuracy = model.evaluate(X_test, y_test)
```

```
Epoch 1/40
7/7 [=====] - 1s 3ms/step - loss: 0.8126 - accuracy: 0.4387
Epoch 2/40
7/7 [=====] - 0s 2ms/step - loss: 0.8098 - accuracy: 0.4481
Epoch 3/40
7/7 [=====] - 0s 3ms/step - loss: 0.7532 - accuracy: 0.4953
Epoch 4/40
7/7 [=====] - 0s 3ms/step - loss: 0.7479 - accuracy: 0.4670
Epoch 5/40
7/7 [=====] - 0s 5ms/step - loss: 0.7442 - accuracy: 0.4906
Epoch 6/40
7/7 [=====] - 0s 2ms/step - loss: 0.7226 - accuracy: 0.5519
Epoch 7/40
7/7 [=====] - 0s 2ms/step - loss: 0.7028 - accuracy: 0.5377
Epoch 8/40
7/7 [=====] - 0s 3ms/step - loss: 0.6868 - accuracy: 0.5660
Epoch 9/40
7/7 [=====] - 0s 3ms/step - loss: 0.6965 - accuracy: 0.5849
Epoch 10/40
7/7 [=====] - 0s 3ms/step - loss: 0.6828 - accuracy: 0.5708
Epoch 11/40
7/7 [=====] - 0s 2ms/step - loss: 0.6582 - accuracy: 0.6321
Epoch 12/40
7/7 [=====] - 0s 2ms/step - loss: 0.6762 - accuracy: 0.6179
Epoch 13/40
7/7 [=====] - 0s 4ms/step - loss: 0.6528 - accuracy: 0.6368
Epoch 14/40
7/7 [=====] - 0s 2ms/step - loss: 0.6356 - accuracy: 0.6887
Epoch 15/40
7/7 [=====] - 0s 2ms/step - loss: 0.6149 - accuracy: 0.6887
Epoch 16/40
7/7 [=====] - 0s 2ms/step - loss: 0.6007 - accuracy: 0.7547
Epoch 17/40
7/7 [=====] - 0s 2ms/step - loss: 0.6115 - accuracy: 0.7123
Epoch 18/40
7/7 [=====] - 0s 3ms/step - loss: 0.5923 - accuracy: 0.6981
Epoch 19/40
7/7 [=====] - 0s 3ms/step - loss: 0.5883 - accuracy: 0.7547
```

Saved successfully!



```

Epoch 20/40
7/7 [=====] - 0s 2ms/step - loss: 0.5880 - accuracy: 0.7264
Epoch 21/40
7/7 [=====] - 0s 2ms/step - loss: 0.5746 - accuracy: 0.7170
Epoch 22/40
7/7 [=====] - 0s 2ms/step - loss: 0.5831 - accuracy: 0.7123
Epoch 23/40
7/7 [=====] - 0s 2ms/step - loss: 0.5704 - accuracy: 0.7642
Epoch 24/40
7/7 [=====] - 0s 2ms/step - loss: 0.5744 - accuracy: 0.7500
Epoch 25/40
7/7 [=====] - 0s 2ms/step - loss: 0.5457 - accuracy: 0.7547
Epoch 26/40
7/7 [=====] - 0s 3ms/step - loss: 0.5360 - accuracy: 0.7783
Epoch 27/40
7/7 [=====] - 0s 2ms/step - loss: 0.5289 - accuracy: 0.7547
Epoch 28/40
7/7 [=====] - 0s 3ms/step - loss: 0.5524 - accuracy: 0.7547
Epoch 29/40

```

#Evaluating models

```

def conf_matrix(matrix,pred):
    class_names= [0,1]# name of classes
    fig, ax = plt.subplots()
    tick_marks = np.arange(len(class_names))
    plt.xticks(tick_marks, class_names)
    plt.yticks(tick_marks, class_names)
    # create heatmap
    sns.heatmap(pd.DataFrame(matrix), annot=True, cmap="YlGnBu" ,fmt='g')
    ax.xaxis.set_label_position("top")
    plt.tight_layout()
    plt.title('Confusion matrix', y=1.1)
    plt.xlabel('Actual label')


```

Saved successfully!

```

#Random Forest
# make class predictions with the model
y_pred = rf.predict(X_test)
cnf_matrix = metrics.confusion_matrix(y_pred,y_test)
conf_matrix(cnf_matrix,y_test)
# calculate prediction
report = classification_report(y_pred,y_test)
print(report)

```




```
#Naive bayes
# make class predictions with the model
y_pred = nb.predict(X_test)
cnf_matrix = metrics.confusion_matrix(y_pred,y_test)
conf_matrix(cnf_matrix,y_test)
# calculate prediction
report = classification_report(y_pred,y_test)
print(report)
```

Saved successfully!



Confusion matrix


Predicted label



```
#svm
# make class predictions with the model
y_pred = svm.predict(X_test)
cnf_matrix = metrics.confusion_matrix(y_pred,y_test)
conf_matrix(cnf_matrix,y_test)
# calculate prediction
report = classification_report(y_pred,y_test)
print(report)
```

Confusion matrix

Predicted label




	precision	recall	f1-score	support
0	0.72	0.81	0.76	42
1	0.73	0.77	0.77	49
macro avg	0.77	0.77	0.77	91
weighted avg	0.77	0.77	0.77	91

Saved successfully!



```
#K-Neighbours
# make class predictions with the model
y_pred = knn.predict(X_test)
cnf_matrix = metrics.confusion_matrix(y_pred,y_test)
conf_matrix(cnf_matrix,y_test)
# calculate prediction
report = classification_report(y_pred,y_test)
print(report)
```



```
#Neural network
# make class predictions with the model
y_predict = np.argmax(model.predict(X_test), axis=-1)
cnf_matrix = metrics.confusion_matrix(y_pred,y_test)
conf_matrix(cnf_matrix,y_test)
report = classification_report(y_pred,y_test)
print(report)
```

Saved successfully!



Confusion matrix

Predicted label



```
from sklearn.metrics import classification_report, confusion_matrix, accuracy_score, fbeta_score
from sklearn.metrics import roc_auc_score
from sklearn.metrics import precision_score
from sklearn.metrics import recall_score
```

```
from sklearn.metrics import f1_score
CM=confusion_matrix(y_test,y_pred_logreg)
sns.heatmap(CM, annot=True)
```

```
TN = CM[0][0]
FN = CM[1][0]
TP = CM[1][1]
FP = CM[0][1]
specificity = TN/(TN+FP)
acc= accuracy_score(y_test, y_pred_logreg)
roc=roc_auc_score(y_test, y_pred_logreg)
prec = precision_score(y_test, y_pred_logreg)
rec = recall_score(y_test, y_pred_logreg)
f1 = f1_score(y_test, y_pred_logreg)
```


```
model_results =pd.DataFrame([['Logistic Regression',acc, prec,rec,specificity, f1,roc]],
                             columns = ['Model', 'Accuracy','Precision', 'Sensitivity','Specificity', 'F1 Score', 'ROC'])
```

```
model_results
```

Saved successfully!



	Precision	Sensitivity	Specificity	F1 Score	ROC
Logistic Regression	0.826087	0.863636	0.829787	0.844444	0.846712



```
from sklearn.metrics import log_loss, roc_auc_score, precision_score, f1_score, recall_score, roc_auc_score
data = {
    'logreg': y_pred_logreg,
```



```

        'nb': y_pred_nb,
        'svm': y_pred_svm,
        'knn': y_pred_knn,
        'rf': y_pred_rf,
    }

models = pd.DataFrame(data)

for column in models:
    CM=confusion_matrix(y_test,models[column])

    TN = CM[0][0]
    FN = CM[1][0]
    TP = CM[1][1]
    FP = CM[0][1]
    specificity = TN/(TN+FP)

    acc= accuracy_score(y_test, models[column])
    roc=roc_auc_score(y_test, models[column])
    prec = precision_score(y_test, models[column])
    rec = recall_score(y_test, models[column])
    f1 = f1_score(y_test, models[column])

    results =pd.DataFrame([[column,acc, prec,rec,specificity, f1,roc]],
        columns = ['Model', 'Accuracy','Precision', 'Sensitivity','Specificity', 'F1 S
    model_results = model_results.append(results, ignore_index = True)

model_results

```

Saved successfully!



			Precision	Sensitivity	Specificity	F1 Score	ROC
			0.826087	0.863636	0.829787	0.844444	0.846712
1	logreg	0.846154	0.826087	0.863636	0.829787	0.844444	0.846712
2	nb	0.802198	0.809524	0.772727	0.829787	0.790698	0.801257
3	svm	0.769231	0.734694	0.818182	0.723404	0.774194	0.770793
4	knn	0.835165	0.808511	0.863636	0.808511	0.835165	0.836074
5	rf	0.857143	0.829787	0.886364	0.829787	0.857143	0.858075

```


def roc_auc_plot(y_true, y_proba, label=' ', l='-', lw=1.0):
    from sklearn.metrics import roc_curve, roc_auc_score
    fpr, tpr, _ = roc_curve(y_true, y_proba[:,1])
    ax.plot(fpr, tpr, linestyle=l, linewidth=lw,
            label="%s (area=%.3f)"%(label,roc_auc_score(y_true, y_proba[:,1])))

```

```
f, ax = plt.subplots(figsize=(12,8))
```

```
roc_auc_plot(y_test,logreg.predict_proba(X_test),label='Logistic Regression ',l='--')
roc_auc_plot(y_test,nb.predict_proba(X_test),label='NB ',l='--')
roc_auc_plot(y_test,svm.predict_proba(X_test),label='SVM ',l='--')
roc_auc_plot(y_test,knn.predict_proba(X_test),label='KNN ',l='--')
roc_auc_plot(y_test,rf.predict_proba(X_test),label='RF ',l='--')
```

```
ax.plot([0,1], [0,1], color='k', linewidth=0.5, linestyle='--',
        )
ax.legend(loc="lower right")
ax.set_xlabel('False Positive Rate')
ax.set_ylabel('True Positive Rate')
ax.set_xlim([0, 1])
ax.set_ylim([0, 1])
ax.set_title('Receiver Operator Characteristic curves')
sns.despine()
```



```
# TODO:- MODEL, TRAINING and TESTING
import pandas as pd
data = pd.read_csv("/content/heart.csv")
print(data.corr())
data.corr()
```

	age	sex	cp	...	ca	thal	target
age	1.000000	-0.098447	-0.068653	...	0.276326	0.068001	-0.225439
sex	-0.098447	1.000000	-0.049353	...	0.118261	0.210041	-0.280937
cp	-0.068653	-0.049353	1.000000	...	-0.181053	-0.161736	0.433798
trestbps	0.279351	-0.056769	0.047608	...	0.101389	0.062210	-0.144931
chol	0.213678	-0.197912	-0.076904	...	0.070511	0.098803	-0.085239
fbs	0.121308	0.045032	0.094444	...	0.137979	-0.032019	-0.028046
restecg	-0.116211	-0.058196	0.044421	...	-0.072042	-0.011981	0.137230
thalach	-0.398522	-0.044020	0.295762	...	-0.213177	-0.096439	0.421741
exang	0.096801	0.141664	-0.394280	...	0.115739	0.206754	-0.436757
oldpeak	0.210013	0.096093	-0.149230	...	0.222682	0.210244	-0.430696
slope	-0.168814	-0.030711	0.119717	...	-0.080155	-0.104764	0.345877
ca	0.276326	0.118261	-0.181053	...	1.000000	0.151832	-0.391724
thal	0.068001	0.210041	-0.161736	...	0.151832	1.000000	-0.344029
target	-0.225439	-0.280937	0.433798	...	-0.391724	-0.344029	1.000000

[14 rows x 14 columns]

	age	sex	cp	trestbps	chol	fbs	restecg	thala
<b>age</b>	1.000000	-0.098447	-0.068653	0.279351	0.213678	0.121308	-0.116211	-0.3985
<b>sex</b>	-0.098447	1.000000	-0.049353	-0.056769	-0.197912	0.045032	-0.058196	-0.0440
<b>cp</b>	-0.068653	-0.049353	1.000000	0.047608	-0.076904	0.094444	0.044421	0.2957
<b>trestbps</b>	0.279351	-0.056769	0.047608	1.000000	0.123174	0.177531	-0.114103	-0.0466
<b>chol</b>	0.213678	-0.197912	-0.076904	0.123174	1.000000	0.013294	-0.151040	-0.0099
<b>fbs</b>	0.121308	0.045032	0.094444	0.177531	0.013294	1.000000	-0.084189	-0.0085
<b>restecg</b>	-0.116211	-0.058196	0.044421	-0.114103	-0.151040	-0.084189	1.000000	0.0441
<b>thalach</b>	-0.398522	-0.044020	0.295762	-0.046698	-0.009940	-0.008567	0.044123	1.0000
			-0.394280	0.067616	0.067023	0.025665	-0.070733	-0.3788
			-0.149230	0.193216	0.053952	0.005747	-0.058770	-0.3441
<b>slope</b>	-0.168814	-0.030711	0.119717	-0.121475	-0.004038	-0.059894	0.093045	0.3867
<b>ca</b>	0.276326	0.118261	-0.181053	0.101389	0.070511	0.137979	-0.072042	-0.2131
<b>thal</b>	0.068001	0.210041	-0.161736	0.062210	0.098803	-0.032019	-0.011981	-0.0964
<b>target</b>	-0.225439	-0.280937	0.433798	-0.144931	-0.085239	-0.028046	0.137230	0.4217

Saved successfully!



```
data=data.rename(columns={'age':'Age','sex':'Sex','cp':'Cp','trestbps':'Trestbps','chol':'Cho
```

```
#New show columns
print(data.columns)
```

```
Index(['Age', 'Sex', 'Cp', 'Trestbps', 'Chol', 'Fbs', 'Restecg', 'Thalach',
      'Exang', 'Oldpeak', 'Slope', 'Ca', 'Thal', 'Target'],
      dtype='object')
```

```
dataX=data.drop('Target',axis=1)
dataY=data['Target']
X_train,X_test,y_train,y_test=train_test_split(dataX,dataY,test_size=0.3,random_state=42)
print('X_train',X_train.shape)
print('X_test',X_test.shape)
print('y_train',y_train.shape)
print('y_test',y_test.shape)
```

```
X_train (212, 13)
X_test (91, 13)
y_train (212,)
y_test (91,)
```

```
#Normalization as the first process
```

```
# Normalize
```

```
X_train=(X_train-np.min(X_train))/(np.max(X_train)-np.min(X_train)).values
X_test=(X_test-np.min(X_test))/(np.max(X_test)-np.min(X_test)).values
```

```
from sklearn.decomposition import PCA
pca=PCA().fit(X_train)
print(pca.explained_variance_ratio_)
print()
print(X_train.columns.values.tolist())
print(pca.components_)
```

```
[0.27983251 0.19244371 0.1190256  0.08185051 0.07970383 0.06329404
 0.02481422 0.01802431 0.01653643]
```

Saved successfully!

```
[ 'Age', 'Sex', 'Cp', 'Trestbps', 'Chol', 'Fbs', 'Restecg', 'Thalach', 'Exang', 'Oldpeak'
[[ 0.06748528  0.46514006 -0.33713815  0.00523708 -0.00675834 -0.02008128
 -0.0181853  -0.17291147  0.72233753  0.17007432 -0.23143612  0.11507873
 0.13122623]
 [ 0.08899783 -0.86322002 -0.24103172  0.05347111  0.06543148 -0.09444788
 0.05938857 -0.11409076  0.37485988  0.03054307 -0.11196353 -0.03895024
 -0.02690692]
 [ 0.14654406 -0.08815103  0.38377419  0.15650123  0.03611168  0.80618106
 -0.16386029 -0.06959297  0.09691224  0.10496249 -0.29659932  0.080185
 -0.03526663]
 [ 0.16278572 -0.00556925 -0.29501656 -0.01030338  0.02152029 -0.15195454
 -0.06535169 -0.1782624  -0.51825582  0.29299049 -0.65122965  0.21392761
 0.04929124]
 [-0.19141487  0.00820035  0.60018286 -0.08955885 -0.07324012 -0.35074511
 0.26120455  0.06601115  0.17659696  0.14948834 -0.47074568 -0.34463381
 -0.01319521]
 [-0.18145061  0.03180462 -0.29522609 -0.02976676 -0.06633548  0.37037866
 0.83843847  0.03970393 -0.10675794 -0.00683364 -0.04770591 -0.11953167
```

```

0.06225589]
[ 0.31042588 -0.02880005  0.36713044  0.09113953  0.02969298 -0.18959162
 0.40349644 -0.20138476  0.03627626  0.14717541  0.23055634  0.66083732
 0.08175907]
[-0.13661882  0.03241732 -0.02414453 -0.3308474 -0.17679525  0.02178087
 0.03093033 -0.19029236  0.0632074 -0.20104409 -0.10414964  0.24216847
 -0.82911761]
[ 0.5272751  0.14295977 -0.00819303  0.41238968  0.01682368 -0.08211305
 0.12443058 -0.41118173 -0.06874053 -0.0907869  0.07947623 -0.48748334
 -0.28793772]
[ 0.07309209 -0.04401741  0.09203714 -0.62824629 -0.11726315  0.09247649
 -0.04833253 -0.60408946 -0.05520955 -0.15508501  0.08138407 -0.16070915
 0.37475708]
[-0.25411673 -0.01547649 -0.01385981 -0.08318313  0.18258015  0.05907082
 -0.05099075 -0.21262928 -0.04363894  0.81460093  0.33386583 -0.15727544
 -0.20434411]
[ 0.54908487  0.0415315 -0.02051409 -0.51363394  0.44670922  0.03592389
 0.08996002  0.4317689  0.04378352  0.08804434 -0.00757107 -0.11116942
 -0.12287669]
[-0.32484899  0.05827329  0.04502966  0.10170275  0.83969328 -0.03101672
 0.03919764 -0.26857643 -0.01542372 -0.29187053 -0.09739615  0.06380959
 -0.00238118]]


```

```

pca = PCA(n_components=8)
pca.fit(X_train)
reduced_data_train = pca.transform(X_train)
#inverse_data = pca.inverse_transform(reduced_data)
plt.scatter(reduced_data_train[:, 0], reduced_data_train[:, 1], label='reduced')
plt.xlabel('First Principal Component')
plt.ylabel('Second Principal Component')
plt.show()

```

Saved successfully!




```

pca = PCA(n_components=8)
pca.fit(X_test)
reduced_data_test = pca.transform(X_test)
#inverse_data = pca.inverse_transform(reduced_data)

```

```
plt.scatter(reduced_data_test[:, 0], reduced_data_test[:, 1], label='reduced')
plt.xlabel('First Principal Component')
plt.ylabel('Second Principal Component')
plt.show()
```



```
reduced_data_train = pd.DataFrame(reduced_data_train, columns=['Dim1', 'Dim2', 'Dim3', 'Dim4', 'Dim5', 'Dim6', 'Dim7', 'Dim8'])
reduced_data_test = pd.DataFrame(reduced_data_test, columns=['Dim1', 'Dim2', 'Dim3', 'Dim4', 'Dim5', 'Dim6', 'Dim7', 'Dim8'])
X_train=reduced_data_train
X_test=reduced_data_test
```

```
def plot_roc_(false_positive_rate,true_positive_rate,roc_auc):
    plt.figure(figsize=(5,5))
    plt.title('Receiver Operating Characteristic')
    plt.plot(false_positive_rate,true_positive_rate, color='red',label = 'AUC = %0.2f' % roc_auc)
    plt.legend(loc='lower right')
    plt.grid(True)
    plt.show()

plt.ylabel('True Positive Rate')
plt.xlabel('False Positive Rate')
plt.show()
```

Saved successfully!

✕ e='--')

```
def plot_feature_importances(gbm):
    n_features = X_train.shape[1]
    plt.barh(range(n_features), gbm.feature_importances_, align='center')
    plt.yticks(np.arange(n_features), X_train.columns)
    plt.xlabel("Feature importance")
    plt.ylabel("Feature")
    plt.ylim(-1, n_features)
    combine_features_list=[
        ('Dim1', 'Dim2', 'Dim3'),
        ('Dim4', 'Dim5', 'Dim5', 'Dim6'),
        ('Dim7', 'Dim8', 'Dim1'),
        ('Dim4', 'Dim8', 'Dim5')
    ]
```

]

```

# TODO :-----Logistic Regression

parameters=[
{
    'penalty':['l1','l2'],
    'C':[0.1,0.4,0.5],
    'random_state':[0]
},
]

for features in combine_features_list:
    print(features)
    print(""*50)

    X_train_set=X_train.loc[:,features]
    X_test_set=X_test.loc[:,features]

    gslog=GridSearchCV(LogisticRegression(),parameters,scoring='accuracy')
    gslog.fit(X_train_set,y_train)
    print('Best parameters set:')
    print(gslog.best_params_)
    print()
    predictions=[
        (gslog.predict(X_train_set),y_train,'Train'),
        (gslog.predict(X_test_set),y_test,'Test'),
    ]

    for pred in predictions:
        print(pred[2] + ' Classification Report:')
        print(""*50)
        print(pred[1],pred[0]))

        print(pred[2] + ' Confusion Matrix:')
        print(confusion_matrix(pred[1], pred[0]))
        print(""*50)

    print(""*50)
    basari=cross_val_score(estimator=LogisticRegression(),X=X_train,y=y_train,cv=12)
    print(basari.mean())
    print(basari.std())
    print(""*50)

    ('Dim1', 'Dim2', 'Dim3')
    *****
    Best parameters set:
    {'C': 0.5, 'penalty': 'l2', 'random_state': 0}

    Train Classification Report:
    *****

```

Saved successfully!



	precision	recall	f1-score	support
0	0.83	0.64	0.72	97
1	0.74	0.89	0.81	115
accuracy			0.77	212
macro avg	0.79	0.76	0.77	212
weighted avg	0.78	0.77	0.77	212

\*\*\*\*\*

Train Confusion Matrix:

```
[[ 62  35]
 [ 13 102]]
```

\*\*\*\*\*

Test Classification Report:

\*\*\*\*\*

	precision	recall	f1-score	support
0	0.74	0.61	0.67	41
1	0.72	0.82	0.77	50
accuracy			0.73	91
macro avg	0.73	0.71	0.72	91
weighted avg	0.73	0.73	0.72	91

\*\*\*\*\*

Test Confusion Matrix:

```
[[25 16]
 [ 9 41]]
```

\*\*\*\*\*

\*\*\*\*\*

0.8259803921568629

0.06685549585135142

\*\*\*\*\*

\*\*\*\*\*

Saved successfully!



```
{'C': 0.4, 'penalty': 'l2', 'random_state': 0}
```

Train Classification Report:

\*\*\*\*\*

	precision	recall	f1-score	support
0	0.66	0.47	0.55	97
1	0.64	0.79	0.71	115
accuracy			0.65	212
macro avg	0.65	0.63	0.63	212
weighted avg	0.65	0.65	0.64	212

\*\*\*\*\*

```
from sklearn.linear_model import LogisticRegression
```


```
lr=LogisticRegression()
lr.fit(X_train,y_train)
```



```
y_pred=lr.predict(X_test)
```

```
y_proba=lr.predict_proba(X_test)
```

```
false_positive_rate, true_positive_rate, thresholds = roc_curve(y_test,y_proba[:,1])
roc_auc = auc(false_positive_rate, true_positive_rate)
plot_roc_(false_positive_rate,true_positive_rate,roc_auc)
```



```
from sklearn.metrics import r2_score,accuracy_score
```


```
print('Error rate :',r2_score(y_test,y_pred))
```

Saved successfully!

```
print('Accuracy rate :',accuracy_score(y_test, y_pred))
print('Logistic TRAIN score with ',format(lr.score(X_train, y_train)))
print('Logistic TEST score with ',format(lr.score(X_test, y_test)))
print()
```

```
Error rate : 0.28975609756097553
Accuracy rate : 0.8241758241758241
Logistic TRAIN score with 0.8537735849056604
Logistic TEST score with 0.8241758241758241
```

```
cm=confusion_matrix(y_test,y_pred)
print(cm)
sns.heatmap(cm,annot=True)
plt.show()
```



# TODO:- K-Nearest Neighbors

```
parameters=[
{
    'n_neighbors':np.arange(2,33),
    'n_jobs':[2,6]
},
]
print("""*50)
for features in combine_features_list:
    print("""*50)

    X_train_set=X_train.loc[:,features]
    X_test_set=X_test.loc[:,features]

    gsknn=GridSearchCV(KNeighborsClassifier(),parameters,scoring='accuracy')
    gsknn.fit(X_train_set,y_train)
    print('Best parameters set:')

    predictions = [
        (gsknn.predict(X_train_set), y_train, 'Train'),
        (gsknn.predict(X_test_set), y_test, 'Test1')
    ]
    for pred in predictions:
        print(pred[2] + ' Classification Report:')
        print("""*50)
        print(classification_report(pred[1], pred[0]))
        print("""*50)
        print(pred[2] + ' Confusion Matrix:')
        print(confusion_matrix(pred[1], pred[0]))
        print("""*50)

    print("""*50)
    basari=cross_val_score(estimator=KNeighborsClassifier(),X=X_train,y=y_train,cv=12)
    print(basari.mean())
```

Saved successfully!



```
print(basari.std())
print("*"*50)
```

```
*****
*****
```

Best parameters set:

```
{'n_jobs': 2, 'n_neighbors': 14}
```

```
*****
```

Train Classification Report:

```
*****
```

	precision	recall	f1-score	support
0	0.75	0.87	0.80	97
1	0.87	0.76	0.81	115
accuracy			0.81	212
macro avg	0.81	0.81	0.81	212
weighted avg	0.82	0.81	0.81	212

```
*****
```

Train Confusion Matrix:

```
[[84 13]
```

```
[28 87]]
```

```
*****
```

Test1 Classification Report:

```
*****
```

	precision	recall	f1-score	support
0	0.74	0.68	0.71	41
1	0.75	0.80	0.78	50
accuracy			0.75	91
macro avg	0.75	0.74	0.74	91
weighted avg	0.75	0.75	0.75	91

Saved successfully!



```
*****
```

```
[[28 13]
```

```
[10 40]]
```

```
*****
```

```
*****
```

```
0.7794117647058824
```

```
0.12371270306248845
```

```
*****
```

```
*****
```

Best parameters set:

```
{'n_jobs': 2, 'n_neighbors': 4}
```

```
*****
```

Train Classification Report:

```
*****
```

	precision	recall	f1-score	support
0	0.70	0.89	0.78	97
1	0.88	0.68	0.76	115
accuracy			0.77	212

macro avg	0.79	0.78	0.77	212
weighted avg	0.80	0.77	0.77	212

\*\*\*\*\*

Train Confusion Matrix:


```
knn=KNeighborsClassifier(n_jobs=2, n_neighbors=22)
knn.fit(X_train,y_train)
```

```
y_pred=knn.predict(X_test)
```

```
y_proba=knn.predict_proba(X_test)
false_positive_rate, true_positive_rate, thresholds = roc_curve(y_test,y_proba[:,1])
roc_auc = auc(false_positive_rate, true_positive_rate)
plot_roc(false_positive_rate,true_positive_rate,roc_auc)
```


```
from sklearn.metrics import r2_score,accuracy_score
```

```
print('Accuracy rate :',accuracy_score(y_test, y_pred))
print("KNN TRAIN score with ",format(knn.score(X_train, y_train)))
print("KNN TEST score with ",format(knn.score(X_test, y_test)))
print()
```




```
Accuracy rate : 0.8241758241758241
KNN TRAIN score with 0.8018867924528302
KNN TEST score with 0.8241758241758241
```

```
cm=confusion_matrix(y_test,y_pred)
print(cm)
sns.heatmap(cm,annot=True)
plt.show()
```



```
n_neighbors = range(1, 17)
train_data_accuracy = []
test1_data_accuracy = []
for n_neigh in n_neighbors:
    knn = KNeighborsClassifier(n_neighbors=n_neigh,n_jobs=5)
    knn.fit(X_train, y_train)
    train_data_accuracy.append(knn.score(X_train, y_train))
    test1_data_accuracy.append(knn.score(X_test, y_test))
plt.plot(n_neighbors, train_data_accuracy, label="Train Data Set")
plt.plot(n_neighbors, test1_data_accuracy, label="Test1 Data Set")
plt.ylabel("Accuracy")
plt.xlabel("Neighbors")
plt.legend()
plt.show()
```




```
n_neighbors = range(1, 17)
k_scores=[]
for n_neigh in n_neighbors:
    knn = KNeighborsClassifier(n_neighbors=n_neigh,n_jobs=5)
    scores=cross_val_score(estimator=knn,X=X_train,y=y_train,cv=12)
```

```
k_scores.append(scores.mean())
print(k_scores)
```

```
[0.7456427015250545, 0.7369281045751634, 0.812091503267974, 0.8071895424836603, 0.779411
```

```
plt.plot(n_neighbors,k_scores)
plt.xlabel('Value of k for KNN')
plt.ylabel("Cross-Validated Accuracy")
plt.show()
```



```
print('Leaf Size :',knn.leaf_size)
print('Metric :',knn.metric_params)
print('Radius :',knn.radius)
print('Weights :',knn.weights)
```

Saved successfully!

```
Metric : None
Radius : None
Weights : uniform
Algorithms : auto
```

```
# DOTO :0 Naive Baes
parameters = [
    {
        'kernel': ['linear'],
        'random_state': [2]
    },
    {
        'kernel': ['rbf'],
        'gamma': [0.9, 0.06, 0.3],
        'random_state': [0],
        'C': [1, 2, 3, 4, 5, 6],
        'degree': [2],
```

```

        'probability':[True]
    },
]

for features in combine_features_list:
    print(""*50)
    X_train_set=X_train.loc[:,features]
    X_test_set=X_test.loc[:,features]

    svc = GridSearchCV(SVC(), parameters,
                        scoring='accuracy')
    svc.fit(X_train_set, y_train)
    print('Best parameters set:')
    print(svc.best_params_)
    print(""*50)
    predictions = [
        (svc.predict(X_train_set), y_train, 'Train'),
        (svc.predict(X_test_set), y_test, 'Test1')
    ]
    for pred in predictions:
        print(pred[2] + ' Classification Report:')
        print(""*50)
        print(classification_report(pred[1], pred[0]))
        print(""*50)
        print(pred[2] + ' Confusion Matrix:')
        print(confusion_matrix(pred[1], pred[0]))
        print(""*50)

    print(""*50)
    basari=cross_val_score(estimator=SVC(),X=X_train,y=y_train,cv=4)
    print(basari.mean())

```

Saved successfully!

\*\*\*\*\*

Best parameters set:

{'C': 6, 'degree': 2, 'gamma': 0.9, 'kernel': 'rbf', 'probability': True, 'random\_state': 0}

\*\*\*\*\*

Train Classification Report:

\*\*\*\*\*

	precision	recall	f1-score	support
0	0.77	0.73	0.75	97
1	0.78	0.82	0.80	115
accuracy			0.78	212
macro avg	0.78	0.77	0.78	212
weighted avg	0.78	0.78	0.78	212

\*\*\*\*\*

Train Confusion Matrix:

[[71 26]

```
[21 94]]
*****
Test1 Classification Report:
*****
              precision    recall  f1-score   support

         0           0.71       0.61       0.66         41
         1           0.71       0.80       0.75         50

 accuracy                   0.71         91
 macro avg           0.71       0.70       0.71         91
 weighted avg        0.71       0.71       0.71         91

*****
Test1 Confusion Matrix:
[[25 16]
 [10 40]]
*****
*****
0.7877358490566038
0.05398831670877171
*****
*****
Best parameters set:
{'C': 2, 'degree': 2, 'gamma': 0.9, 'kernel': 'rbf', 'probability': True, 'random_state': 0}
*****
Train Classification Report:
*****
              precision    recall  f1-score   support

         0           0.64       0.51       0.57         97
         1           0.65       0.77       0.70        115

 accuracy                   0.65        212
 macro avg           0.64       0.63       0.63        212
 weighted avg        0.65       0.64       0.64        212

*****
Train Confusion Matrix:
```

Saved successfully!


```
from sklearn import svm
```

```
svc=svm.SVC(C=5,degree=2,gamma=0.06,kernel='rbf',probability=True,random_state=0)
svc.fit(X_train,y_train)
```

```
y_pred=svc.predict(X_test)
```

```
y_proba=svc.predict_proba(X_test)
false_positive_rate, true_positive_rate, thresholds = roc_curve(y_test,y_proba[:,1])
roc_auc = auc(false_positive_rate, true_positive_rate)
plot_roc_(false_positive_rate,true_positive_rate,roc_auc)
```





```
from sklearn.metrics import r2_score, accuracy_score
```


```
print('Accuracy rate :', accuracy_score(y_test, y_pred))
print("SVC TRAIN score with ", format(svc.score(X_train, y_train)))
print("SVC TEST score with ", format(svc.score(X_test, y_test)))
print()
```

```
Accuracy rate : 0.8351648351648352
SVC TRAIN score with 0.8301886792452831
SVC TEST score with 0.8351648351648352
```

```
cm=confusion_matrix(y_test,y_pred)
print(cm)
```

Saved successfully!

```
[[33  8]
 [ 7 43]]
```



```

# Random Forest
parameters = [
    {
        'max_depth': np.arange(1, 10),
        'min_samples_split': np.arange(2, 5),
        'random_state': [3],
        'n_estimators': np.arange(10, 20)
    },
]

for features in combine_features_list:
    print(""*50)

    X_train_set=X_train.loc[:,features]
    X_test1_set=X_test.loc[:,features]

    tree=GridSearchCV(RandomForestClassifier(),parameters,scoring='accuracy')
    tree.fit(X_train_set, y_train)

    print('Best parameters set:')
    print(tree.best_params_)
    print(""*50)
    predictions = [
        (tree.predict(X_train_set), y_train, 'Train'),
        (tree.predict(X_test1_set), y_test, 'Test1')
    ]

    for pred in predictions:

        print(pred[2] + ' Classification Report:')
        print(""*50)
        print(pred[2] + ' Confusion Matrix:')
        print(pred[1], pred[0]))
        print(confusion_matrix(pred[1], pred[0]))
        print(""*50)

    print(""*50)
    basari=cross_val_score(estimator=RandomForestClassifier(),X=X_train,y=y_train,cv=4)
    print(basari.mean())
    print(basari.std())
    print(""*50)

    *****
    Best parameters set:
    {'max_depth': 6, 'min_samples_split': 3, 'n_estimators': 18, 'random_state': 3}
    *****
    Train Classification Report:
    *****
    precision    recall  f1-score   support

    0             0.92      0.92      0.92         97

```

Saved successfully!



```

1          0.93      0.93      0.93      115

accuracy          0.92      212
macro avg        0.92      0.92      0.92      212
weighted avg     0.92      0.92      0.92      212

```

```
*****
```

```
Train Confusion Matrix:
```

```
[[ 89   8]
 [   8 107]]
```

```
*****
```

```
Test1 Classification Report:
```

```
*****
```

```

              precision    recall  f1-score   support

0           0.74         0.68         0.71         41
1           0.75         0.80         0.78         50

accuracy          0.75         0.75         0.75         91
macro avg         0.75         0.74         0.74         91
weighted avg      0.75         0.75         0.75         91

```

```
*****
```

```
Test1 Confusion Matrix:
```

```
[[28 13]
 [10 40]]
```

```
*****
```

```
*****
```

```
0.7830188679245284
```

```
0.021094980919809332
```

```
*****
```

```
*****
```

```
Best parameters set:
```

```
{'max_depth': 9, 'min_samples_split': 3, 'n_estimators': 13, 'random_state': 3}
```

```
*****
```

Saved successfully!



```

              precision    recall  f1-score   support

0           0.94         0.99         0.96         97
1           0.99         0.95         0.97        115

accuracy          0.97         0.97         0.97        212
macro avg         0.97         0.97         0.97        212
weighted avg      0.97         0.97         0.97        212

```

```
*****
```

```
Train Confusion Matrix:
```

```
[[ 96   1]
```


```
rfc=RandomForestClassifier(max_depth=7,min_samples_split=4,n_estimators=19,random_state=3)
rfc.fit(X_train,y_train)
```

```
y_pred=rfc.predict(X_test)
```

```

y_proba=rfc.predict_proba(X_test)
false_positive_rate, true_positive_rate, thresholds = roc_curve(y_test,y_proba[:,1])
roc_auc = auc(false_positive_rate, true_positive_rate)
plot_roc_(false_positive_rate,true_positive_rate,roc_auc)

```



```

from sklearn.metrics import r2_score,accuracy_score
print('Accuracy rate :',accuracy_score(y_test, y_pred))
print("RandomForestClassifier TRAIN score with ",format(rfc.score(X_train, y_train)))
print("RandomForestClassifier TEST score with ",format(rfc.score(X_test, y_test)))
print()

```

Accuracy rate : 0.8241758241758241

RandomForestClassifier TRAIN score with 0.9764150943396226

RandomForestClassifier TEST score with 0.8241758241758241

Saved successfully!

```

cm=confusion_matrix(y_test,y_pred)
print(cm)
sns.heatmap(cm,annot=True)
plt.show()

```

```
[[33  8]
 [ 8 42]]
```



```
for i in range(1,10):
    rf = RandomForestClassifier(n_estimators=i, random_state = 3, max_depth=7)
    rf.fit(X_train, y_train)
    print("TEST set score w/ " +str(i)+" estimators: {:.5}".format(rf.score(X_test, y_test)))
```

```
TEST set score w/ 1 estimators: 0.74725
TEST set score w/ 2 estimators: 0.7033
TEST set score w/ 3 estimators: 0.75824
TEST set score w/ 4 estimators: 0.78022
TEST set score w/ 5 estimators: 0.78022
TEST set score w/ 6 estimators: 0.83516
TEST set score w/ 7 estimators: 0.85714
TEST set score w/ 8 estimators: 0.83516
TEST set score w/ 9 estimators: 0.85714
```

```
# TODO:- SVM
```

```
parameters = [
{
    'random_state': [42],
    },
]
for features in combine_features_list:
    print(""*50)
    X_train_set=X_train.loc[:,features]
    X_test1_set=X_test.loc[:,features]
```

Saved successfully!



```
ters, scoring='accuracy')
```

```
dtr.fit(X_train_set, y_train)
print('Best parameters set:')
print(dtr.best_params_)
print(""*50)
predictions = [
(dtr.predict(X_train_set), y_train, 'Train'),
(dtr.predict(X_test1_set), y_test, 'Test1')
]
for pred in predictions:
    print(pred[2] + ' Classification Report:')
    print(""*50)
    print(classification_report(pred[1], pred[0]))
    print(""*50)
    print(pred[2] + ' Confusion Matrix:')
    print(confusion_matrix(pred[1], pred[0]))
    print(""*50)
```

```
print(""*50)
```

```
basari=cross_val_score(estimator=SVC(),X=X_train,y=y_train,cv=4)
print(basari.mean())
print(basari.std())
print(" "*50)
```

```
*****
```

Best parameters set:

```
{'random_state': 42}
```

```
*****
```

Train Classification Report:

```
*****
```

	precision	recall	f1-score	support
0	0.78	0.69	0.73	97
1	0.76	0.83	0.80	115
accuracy			0.77	212
macro avg	0.77	0.76	0.76	212
weighted avg	0.77	0.77	0.77	212

```
*****
```

Train Confusion Matrix:

```
[[67 30]
 [19 96]]
```

```
*****
```

Test1 Classification Report:

```
*****
```

	precision	recall	f1-score	support
0	0.74	0.61	0.67	41
1	0.72	0.82	0.77	50
accuracy			0.73	91
			0.72	91
			0.72	91

Saved successfully!

```
*****
```

Test1 Confusion Matrix:

```
[[25 16]
 [ 9 41]]
```

```
*****
```

```
*****
```

```
0.7877358490566038
```

```
0.05398831670877171
```

```
*****
```

```
*****
```

Best parameters set:

```
{'random_state': 42}
```

```
*****
```

Train Classification Report:

```
*****
```

	precision	recall	f1-score	support
0	0.68	0.55	0.61	97

1	0.67	0.78	0.72	115
accuracy			0.67	212
macro avg	0.68	0.66	0.66	212
weighted avg	0.68	0.67	0.67	212

\*\*\*\*\*

Train Confusion Matrix:

-- -- --

```
from sklearn import svm
```

```
SV=svm.SVC(random_state = 1,probability=True)
```

```
SV.fit(X_train,y_train)
```


```
y_pred=SV.predict(X_test)
```

```
y_proba=SV.predict_proba(X_test)
```

```
false_positive_rate, true_positive_rate, thresholds = roc_curve(y_test,y_proba[:,1])
```

```
roc_auc = auc(false_positive_rate, true_positive_rate)
```

```
plot_roc(false_positive_rate,true_positive_rate,roc_auc)
```



```
from sklearn.neural_network import MLPClassifier
```

```
MLP = MLPClassifier(hidden_layer_sizes=(32), learning_rate_init=0.001, max_iter=150)
```

```
model = MLP.fit(X_train, y_train)
```

```
MLP_predict = MLP.predict(X_test)
```

```
MLP_conf_matrix = confusion_matrix(y_test, MLP_predict)
```

```
MLP_acc_score = accuracy_score(y_test, MLP_predict)
```

```
#Printing the confusion matrix and accuracy scoresprint("confussion matrix")
```

```
print(MLP_conf_matrix)
print("\n")
print(classification_report(y_test,MLP_predict))
print("Accuracy of Multilayer Perceptron classifier: {:.3f}".format(MLP_acc_score*100), '%\n')
```

```
[[34  7]
 [ 8 42]]
```

	precision	recall	f1-score	support
0	0.81	0.83	0.82	41
1	0.86	0.84	0.85	50
accuracy			0.84	91
macro avg	0.83	0.83	0.83	91
weighted avg	0.84	0.84	0.84	91

Accuracy of Multilayer Perceptron classifier: 83.516 %

Saved successfully!

